

RankAlign Training Dynamics Analysis

When Does RankAlign / TC Improve?

Automated Analysis Report

June 9, 2026

Abstract

This report analyzes training dynamics for RankAlign models across four model/task combinations (gemma-2-9b-it and Qwen3.5-9B on membership and IFEval tasks). We characterize **(a)** when RankAlign improves correlation over the base model, **(b)** when the full method improves over basic RankAlign, and **(c)** when typicality correction (TC) specifically improves over non-TC training. We also report training diagnostics from WandB logs and additional statistics (concordance, score delta distributions).

Contents

1	Q1: When Does Each Component Help?	2
1.1	Main Comparison: Gen ROC AUC (TC, Self-Scoring)	2
1.2	Key Findings	2
1.3	Visualizations	2
2	Q2: Training Diagnostics	3
2.1	Loss Curves from WandB	3
2.2	Loss Component Breakdown	5
2.3	Score Explosion Diagnostic	6
3	Q3: Additional Statistics	7
3.1	Generator-Validator Concordance	7
3.2	Spearman Correlation (Generator vs Validator)	7
3.3	Score Delta Distributions	8
3.4	Training Dynamics	8
3.5	Generator-Validator Scatter Plots	9
3.6	Per-Category Analysis	10
4	Focused TC Comparison	10
4.1	Self-TC vs No-TC (s3 vs s4)	10
4.2	Neg-TC vs No-TC (s7 vs s4)	10
5	Conclusions	11

1 Q1: When Does Each Component Help?

1.1 Main Comparison: Gen ROC AUC (TC, Self-Scoring)

Table 1 shows the generator ROC AUC (typicality-corrected) for the final model (epoch 2) evaluated on the training set with self-TC scoring.

Table 1: Generator ROC AUC (TC, self-scoring) at epoch 2. Bold = best per row. Δ relative to base.

Model / Task	Base	s1 (SFT)	s2 (RA basic)	s4 (RA full)	s3 (TC-self)
Gemma / Membership	0.851	—	0.941	0.948	0.976
Gemma / IFEval	0.670	0.586	0.482	0.753	0.801
Qwen / Membership	0.799	0.902	0.880	0.899	0.969
Qwen / IFEval	0.606	0.599	0.558	0.711	0.743

1.2 Key Findings

1. **s3 (TC-self) is the best setting in every single combination.**
2. **Basic RankAlign (s2) hurts on IFEval** — Gemma IFEval drops from 0.670 to 0.482 (below chance). Without NLL constraints, generator scores explode.
3. **The full method (s4) consistently helps**, adding +0.08–0.11 over base across all combos.
4. **TC-self (s3) adds +0.03–0.07 over s4** — a consistent, meaningful improvement from typicality-aware training.

1.3 Visualizations




Figure 1: Generator ROC AUC across settings and model/task combinations (epoch 2, self-TC scoring, train set).




Figure 2: Improvement (Δ) in gen ROC AUC over base model.

2 Q2: Training Diagnostics

2.1 Loss Curves from WandB

Training was monitored via WandB for gemma-ifeval and qwen-membership/ifeval (gemma-membership was trained pre-WandB).




Figure 3: Training loss curves for gemma-2-9b-it on IFEval.




Figure 4: Training loss curves for Qwen3.5-9B on membership.




Figure 5: Training loss curves for Qwen3.5-9B on IFEval.

2.2 Loss Component Breakdown

Table 2 reveals the root cause of s2’s failure on IFEval: without an NLL-G constraint, the generator’s NLL explodes to **3120** (vs. 3.8–4.3 for constrained settings).

Table 2: Final-third loss components (Gemma IFEval). s2’s NLL-G explosion explains its poor gen_roc.

Setting	Total	Pref	NLL-G	NLL-V
s1 (SFT)	5.99	0.00	5.99	0.000
s2 (basic)	11.56	11.56	3119.8	0.000
s3 (TC-self)	6.40	2.64	3.76	0.001
s4 (full)	6.15	1.81	4.33	0.009
s7 (TC-neg)	5.13	1.37	3.77	0.001

Key insight: On the simpler membership task (Qwen), s2 keeps NLL-G reasonable (2.42) without explicit constraint. The constraint is only critical for complex tasks where the optimization landscape allows score explosion.




Figure 6: Score spread evolution during training. s2 on Gemma IFEval shows unbounded generator score growth.




Figure 7: Preference loss evolution. All constrained settings converge; s2 achieves low preference loss at the cost of NLL explosion.

2.3 Score Explosion Diagnostic

A key diagnostic is the mean pairwise absolute difference in generator scores (`gen_delta.mean`). Large values indicate score explosion during training.

Table 3: Generator score spread diagnostic. Values >1000 indicate training instability.

Task	Setting	gen_delta_mean	gen_roc	Status
Gemma IFEval	s2 (basic)	13,421	0.482	Exploded
Gemma IFEval	s4 (full)	691	0.753	OK
Gemma IFEval	s3 (TC-self)	216	0.801	Best
Qwen IFEval	s2 (basic)	1,031	0.558	Borderline
Qwen IFEval	s4 (full)	408	0.711	OK
Qwen IFEval	s3 (TC-self)	139	0.743	Best
Gemma Membership	s2	9.1	0.941	OK
Gemma Membership	s4	10.9	0.948	OK
Gemma Membership	s3	10.7	0.976	Best

Pattern: TC-self training produces the most compact score distributions (lowest gen_delta_mean) and the best gen_roc. The NLL constraints prevent explosion, and TC further tightens the score distribution.

3 Q3: Additional Statistics

3.1 Generator-Validator Concordance

Concordance measures the fraction of item pairs where the generator and validator agree on the ordering direction.

Table 4: Concordance at epoch 2 (self-TC scoring). Chance = 0.5.

Model / Task	Base	s1	s2	s4	s3
Gemma / Membership	0.705	—	0.794	0.787	0.803
Gemma / IFEval	0.548	0.561	0.370	0.677	0.694
Qwen / Membership	0.671	0.734	0.759	0.752	0.807
Qwen / IFEval	0.545	0.553	0.571	0.641	0.646

Note: s2 on Gemma IFEval has concordance 0.370 (below chance), confirming the generator scores are *anti-correlated* with the validator.

3.2 Spearman Correlation (Generator vs Validator)

Table 5: Spearman ρ between generator and validator scores (epoch 2, self-TC).

Model / Task	Base	s1	s2	s4	s3
Gemma / Membership	0.585	—	0.798	0.795	0.826
Gemma / IFEval	0.141	0.178	−0.374	0.504	0.567
Qwen / Membership	0.496	0.665	0.723	0.715	0.830
Qwen / IFEval	0.119	0.151	0.205	0.404	0.431

3.3 Score Delta Distributions




Figure 8: Pairwise score delta distributions for gemma membership. Generator (blue) and validator (orange) distributions. Tighter = better calibrated.

3.4 Training Dynamics




Figure 9: Training dynamics for gemma membership: Gen ROC and Spearman across epochs.




Figure 10: Concordance across all model/task combinations.

3.5 Generator-Validator Scatter Plots




Figure 11: Generator vs validator scores for Base, s4, and s3 (self-TC scoring, gemma membership). Training improves the separation between positive (green) and negative (red) items in both generator and validator space.

3.6 Per-Category Analysis




Figure 12: Left: distribution of per-category improvements (Base $\rightarrow$ s3). Right: improvements from TC-self specifically (s4 $\rightarrow$ s3). Both are consistently positive.

Categories where the base model was weakest (gen_roc $\approx$ 0.5–0.7) benefit most from training. For example, “medical specialty” improves from 0.504 to 1.000 (+0.496). Categories already near ceiling show minimal gains (expected).

4 Focused TC Comparison

4.1 Self-TC vs No-TC (s3 vs s4)

When evaluated with self-TC scoring, the typicality-corrected models consistently outperform the non-TC baseline (s4):

- s3 (TC-self, full): +0.028 over s4
- s6 (TC-self fsx, low δ): +0.023 over s4
- s11 (TC-self vlo, no fsx/ppd): +0.022 over s4
- s5 (TC-self basic): +0.010 over s4

4.2 Neg-TC vs No-TC (s7 vs s4)

When evaluated with neg-TC scoring, the improvement from TC-neg training is even larger:

- s12 (TC-neg vlo): +0.075 over s4
- s7 (TC-neg, full): +0.072 over s4

This suggests that TC-trained models learn better internal representations that are specifically captured by the matching TC scoring method.




Figure 13: TC comparison: self-TC variants (left) and neg-TC variants (right) against s4 baseline.




Figure 14: Epoch-by-epoch dynamics for TC comparison. TC-trained models converge quickly and maintain advantage throughout training.

5 Conclusions

1. **RankAlign works when the full method (s4) is used** — the NLL-V/G constraints are essential for preventing score explosion, especially on complex tasks like IFEval.
2. **TC-self training (s3) consistently provides the best results** across all metrics (gen ROC, spearman, concordance) and all model/task combinations.
3. **Basic RankAlign (s2) is unreliable for complex tasks** — it can actually *hurt* performance (IFEval). The full method is needed.
4. **Score spread (gen_delta_mean) is a powerful training diagnostic** — values >1000 indicate instability. TC-self produces the tightest distributions.
5. **Concordance is a complementary metric to ROC AUC** — it directly measures the alignment between generator and validator ranking functions.