

RankAlign Training Dynamics Analysis

When Does RankAlign / TC Improve?

Automated Analysis Report

June 9, 2026

Abstract

This report analyzes training dynamics for RankAlign models across four model/task combinations (gemma-2-9b-it and Qwen3.5-9B on membership and IFEval tasks). We characterize **(a)** when RankAlign improves correlation over the base model, **(b)** when the full method improves over basic RankAlign, and **(c)** when typicality correction (TC) specifically improves over non-TC training. We also report training diagnostics from WandB logs and additional statistics (concordance, score delta distributions).

Contents

1	Q1: When Does Each Component Help?	2
1.1	Main Comparison: Gen ROC AUC (TC, Self-Scoring)	2
1.2	Key Findings	2
1.3	Visualizations	2
2	Q2: Training Diagnostics	3
2.1	Loss Curves from WandB	3
2.1.1	IFEval	3
2.1.2	Membership	5
2.2	Loss Component Breakdown	7
2.3	Score Explosion Diagnostic	8
3	Q3: Additional Statistics	8
3.1	Generator-Validator Concordance	8
3.2	Spearman Correlation (Generator vs Validator)	9
3.3	Score Delta Distributions	9
3.4	Training Dynamics	10
3.5	Generator-Validator Scatter Plots	11
3.6	Per-Category Analysis	11
4	Focused TC Comparison	11
4.1	Self-TC vs No-TC (s4 vs s3)	11
4.2	Neg-TC vs No-TC (s7 vs s3)	12
5	Train vs Test Generalization	12
6	Conclusions	13

1 Q1: When Does Each Component Help?

1.1 Main Comparison: Gen ROC AUC (TC, Self-Scoring)

Table 1 shows the generator ROC AUC (typicality-corrected) for the final model (epoch 2) evaluated on the training set with self-TC scoring.

Table 1: Generator ROC AUC (TC, self-scoring) at epoch 2. Bold = best per row. Δ relative to base.

Model / Task	Base	s1 (SFT)	s2 (RA basic)	s3 (RA full)	s4 (TC-self)
Gemma / Membership	0.851	—	0.941	0.948	0.976
Gemma / IFEval	0.670	0.586	0.482	0.753	0.801
Qwen / Membership	0.799	0.902	0.880	0.899	0.969
Qwen / IFEval	0.606	0.599	0.558	0.711	0.743

1.2 Key Findings

1. **s4 (TC-self) is the best setting in every single combination.**
2. **Basic RankAlign (s2) hurts on IFEval** — Gemma IFEval drops from 0.670 to 0.482 (below chance). Without NLL constraints, generator scores explode.
3. **The full method (s3) consistently helps**, adding +0.08–0.11 over base across all combos.
4. **TC-self (s4) adds +0.03–0.07 over s3** — a consistent, meaningful improvement from typicality-aware training.

1.3 Visualizations




Figure 1: Generator ROC AUC across settings and model/task combinations (epoch 2, self-TC scoring, train set).

2 Q2: Training Diagnostics

2.1 Loss Curves from WandB

Training was monitored via WandB for gemma-ifeval, qwen-membership, and qwen-ifeval. Gemma-membership was trained before WandB logging was added.

Smoothing. All loss curves in this section are smoothed using a running average with a window of 20 steps. To replicate in WandB: open a line plot panel’s settings → Data tab → Smoothing Type: “Running Average” → Smoothing Parameter: 20. Raw (unsmoothed) curves are available directly on WandB.

2.1.1 IFEval




Figure 2: Loss curves: SFT vs RankAlign vs Ours (gemma-2-9b-it, IFEval).

Ours variants (TC comparison) — gemma-9b-it ifeval




Figure 3: Loss curves: Ours TC comparison (gemma-2-9b-it, IFEval).

SFT vs RankAlign vs Ours — qwen3.5-9b ifeval




Figure 4: Loss curves: SFT vs RankAlign vs Ours (Qwen3.5-9B, IFEval).

Ours variants (TC comparison) — qwen3.5-9b ifeval




Figure 5: Loss curves: Ours TC comparison (Qwen3.5-9B, IFEval).

2.1.2 Membership

Note: gemma-2-9b-it membership was trained before WandB logging was added; no loss curves are available for that combination.

SFT vs RankAlign vs Ours — qwen3.5-9b membership




Figure 6: Loss curves: SFT vs RankAlign vs Ours (Qwen3.5-9B, membership).

Ours variants (TC comparison) — qwen3.5-9b membership




Figure 7: Loss curves: Ours TC comparison (Qwen3.5-9B, membership).

2.2 Loss Component Breakdown

Table 2 reveals the root cause of s2’s failure on IFEval: without an NLL-G constraint, the generator’s NLL explodes to **3120** (vs. 3.8–4.3 for constrained settings).

Table 2: Final-third loss components (Gemma IFEval). s2’s NLL-G explosion explains its poor gen.roc.

Setting	Total	Pref	NLL-G	NLL-V
s1 (SFT)	5.99	0.00	5.99	0.000
s2 (basic)	11.56	11.56	3119.8	0.000
s4 (TC-self)	6.40	2.64	3.76	0.001
s3 (full)	6.15	1.81	4.33	0.009
s7 (TC-neg)	5.13	1.37	3.77	0.001

Key insight: On the simpler membership task (Qwen), s2 keeps NLL-G reasonable (2.42) without explicit constraint. The constraint is only critical for complex tasks where the optimization landscape allows score explosion.




Figure 8: Score spread evolution during training. s2 on Gemma IFEval shows unbounded generator score growth.




Figure 9: Preference loss evolution. All constrained settings converge; s2 achieves low preference loss at the cost of NLL explosion.

2.3 Score Explosion Diagnostic

A key diagnostic is the mean pairwise absolute difference in generator scores (`gen_delta_mean`). Large values indicate score explosion during training.

Table 3: Generator score spread diagnostic. Values >1000 indicate training instability.

Task	Setting	<code>gen_delta_mean</code>	<code>gen_roc</code>	Status
Gemma IFEval	s2 (basic)	13,421	0.482	Exploded
Gemma IFEval	s3 (full)	691	0.753	OK
Gemma IFEval	s4 (TC-self)	216	0.801	Best
Qwen IFEval	s2 (basic)	1,031	0.558	Borderline
Qwen IFEval	s3 (full)	408	0.711	OK
Qwen IFEval	s4 (TC-self)	139	0.743	Best
Gemma Membership	s2	9.1	0.941	OK
Gemma Membership	s3	10.9	0.948	OK
Gemma Membership	s4	10.7	0.976	Best

Pattern: TC-self training produces the most compact score distributions (lowest `gen_delta_mean`) and the best `gen_roc`. The NLL constraints prevent explosion, and TC further tightens the score distribution.

3 Q3: Additional Statistics

3.1 Generator-Validator Concordance

Concordance measures the fraction of item pairs where the generator and validator agree on the ordering direction.

Table 4: Concordance at epoch 2 (self-TC scoring). Chance = 0.5.

Model / Task	Base	s1	s2	s3	s4
Gemma / Membership	0.705	—	0.794	0.787	0.803
Gemma / IFEval	0.548	0.561	0.370	0.677	0.694
Qwen / Membership	0.671	0.734	0.759	0.752	0.807
Qwen / IFEval	0.545	0.553	0.571	0.641	0.646

Note: s2 on Gemma IFEval has concordance 0.370 (below chance), confirming the generator scores are *anti-correlated* with the validator.

3.2 Spearman Correlation (Generator vs Validator)

Table 5: Spearman ρ between generator and validator scores (epoch 2, self-TC).

Model / Task	Base	s1	s2	s3	s4
Gemma / Membership	0.585	—	0.798	0.795	0.826
Gemma / IFEval	0.141	0.178	−0.374	0.504	0.567
Qwen / Membership	0.496	0.665	0.723	0.715	0.830
Qwen / IFEval	0.119	0.151	0.205	0.404	0.431

3.3 Score Delta Distributions




Figure 10: Pairwise score delta distributions for gemma membership. Generator (blue) and validator (orange) distributions. Tighter = better calibrated.

3.4 Training Dynamics




Figure 11: Training dynamics for gemma membership: Gen ROC and Spearman across epochs.




Figure 12: Concordance across all model/task combinations.

3.5 Generator-Validator Scatter Plots




Figure 13: Generator vs validator scores for Base, s3, and s4 (self-TC scoring, gemma membership). Training improves the separation between positive (green) and negative (red) items in both generator and validator space.

3.6 Per-Category Analysis




Figure 14: Left: distribution of per-category improvements (Base $\rightarrow$ s4). Right: improvements from TC-self specifically (s3 $\rightarrow$ s4). Both are consistently positive.

Categories where the base model was weakest ($\text{gen_roc} \approx 0.5\text{--}0.7$) benefit most from training. For example, “medical specialty” improves from 0.504 to 1.000 (+0.496). Categories already near ceiling show minimal gains (expected).

4 Focused TC Comparison

4.1 Self-TC vs No-TC (s4 vs s3)

When evaluated with self-TC scoring, the typicality-corrected models consistently outperform the non-TC baseline (s3):

- s4 (TC-self, full): +0.028 over s3
- s6 (TC-self fsx, low δ): +0.023 over s3
- s11 (TC-self vlo, no fsx/ppd): +0.022 over s3

- s5 (TC-self basic): +0.010 over s3

4.2 Neg-TC vs No-TC (s7 vs s3)

When evaluated with neg-TC scoring, the improvement from TC-neg training is even larger:

- s12 (TC-neg v1o): +0.075 over s3
- s7 (TC-neg, full): +0.072 over s3

This suggests that TC-trained models learn better internal representations that are specifically captured by the matching TC scoring method.




Figure 15: TC comparison: self-TC variants (left) and neg-TC variants (right) against s3 baseline.




Figure 16: Epoch-by-epoch dynamics for TC comparison. TC-trained models converge quickly and maintain advantage throughout training.

5 Train vs Test Generalization

A critical finding emerges when comparing train-set dynamics with held-out test sets.

Table 6: Train vs test gen_roc: gemma-2-9b-it membership/rosch (basetyp scoring).

Setting	Train (membership)	Test (rosch)	
s11 (TC-self vlo)	0.970	0.924	−0.046
s4 (TC-self full)	0.976	0.920	−0.056
s5 (TC-self basic)	0.958	0.913	−0.045
s2 (RA basic)	0.941	0.891	−0.050
s3 (RA full)	0.948	0.885	−0.063

Key finding: s3 (full method, no TC) outperforms s2 on training categories but *underperforms* on test categories (different Rosch categories). Meanwhile, all TC-trained models (s4, s11, s5) maintain their advantage on both splits. This suggests **TC promotes generalization** — it prevents overfitting to training-specific patterns.




Figure 17: Train vs test set gen_roc. Left: membership train $\leftrightarrow$ rosch test. Right: ifeval train $\leftrightarrow$ ifeval OOD test.

6 Conclusions

1. **RankAlign works when the full method (s3) is used** — the NLL-V/G constraints are essential for preventing score explosion, especially on complex tasks like IFEval.
2. **TC-self training (s4) consistently provides the best results** across all metrics (gen ROC, spearman, concordance) and all model/task combinations.
3. **TC promotes generalization:** TC-trained models maintain their advantage on held-out test categories, while non-TC models (s3) can overfit to training patterns.
4. **Basic RankAlign (s2) is unreliable for complex tasks** — it can actually *hurt* performance (IFEval). The full method is needed.
5. **Score spread (gen_delta_mean) is a powerful training diagnostic** — values >1000 indicate instability. TC-self produces the tightest distributions.
6. **Concordance is a complementary metric to ROC AUC** — it directly measures the alignment between generator and validator ranking functions.