

Task/Method	eval-TC	ROC _G			ρ		
		orig	rerun	Δ	orig	rerun	Δ
IFEval/SFT	self	67.3	66.2	+1.1	22.9	20.6	+2.2
	self-base	61.7	62.5	-0.9	14.3	13.0	+1.3
	neg	46.7	48.8	-2.1	6.9	3.2	+3.6
	neg-base	50.0	49.9	+0.1	-0.2	-3.6	+3.4
IFEval/RankAlign	self	83.2	67.9	+15.3	62.5	26.3	+36.2
	self-base	81.9	58.3	+23.5	62.0	-1.3	+63.3
	neg	49.6	57.7	-8.1	-0.0	5.0	-5.1
	neg-base	77.3	53.1	+24.2	59.2	-1.6	+60.8
Hyp/SFT	self	80.2	82.6	-2.4	57.4	61.3	-3.9
	self-base	79.6	81.3	-1.7	56.1	59.1	-3.0
	neg	80.7	84.2	-3.5	55.2	59.7	-4.5
	neg-base	79.9	80.7	-0.9	55.5	56.2	-0.7
Hyp/FLORA-Neg	neg	85.6	—	—	63.1	—	—
	neg-base	88.3	86.7	+1.6	71.7	70.2	+1.5

Table 1: **Original vs wandb-rerun** (Qwen3.5-9B): two independently-trained checkpoints of the same setting, evaluated identically. *orig* = paper checkpoint recovered from HF (latkes); *rerun* = the mll wandb rerun. Δ = orig – rerun (pp). IFEval is OOD (20 prompts); Hyponymy is the 10 rosch tasks. Only the four cells with both checkpoints are shown (qwen ifeval s1/s2, rosch s1/s7). All epoch 2.