

Method	eval	Hyponymy				IFEval			
		G2-9b-it		Q3.5-9b		G2-9b-it		Q3.5-9b	
		ROC_G	ρ	ROC_G	ρ	ROC_G	ρ	ROC_G	ρ
Base	self	77.7 $\pm$ 2.2	55.3 $\pm$ 2.7	73.6 $\pm$ 2.2	45.6 $\pm$ 2.4	78.3 $\pm$ 2.2 [†]	38.9 $\pm$ 5.6 [†]	69.3 $\pm$ 3.7 [†]	24.7 $\pm$ 7.0 [†]
Base	neg	81.2 $\pm$ 4.0	52.6 $\pm$ 6.2	73.8 $\pm$ 2.7	43.9 $\pm$ 3.1	59.3 $\pm$ 2.7 [†]	23.9 $\pm$ 4.7 [†]	46.2 $\pm$ 3.8 [†]	-7.6 $\pm$ 5.8 [†]
SFT	self	85.4 $\pm$ 1.4	62.2 $\pm$ 2.4	77.3 $\pm$ 2.0	52.7 $\pm$ 2.7	58.2 $\pm$ 4.0 [†]	13.9 $\pm$ 7.3 [†]	62.5 $\pm$ 4.1 [†]	13.0 $\pm$ 6.2 [†]
SFT	neg	83.5 $\pm$ 1.1	54.6 $\pm$ 2.7	75.6 $\pm$ 2.2	47.6 $\pm$ 3.8	49.3 $\pm$ 3.8 [†]	4.6 $\pm$ 6.7 [†]	49.9 $\pm$ 4.3 [†]	-3.6 $\pm$ 5.8 [†]
RankAlign	self	89.1 $\pm$ 1.2	68.8 $\pm$ 2.4	80.6 $\pm$ 2.0	58.9 $\pm$ 2.9	43.4 $\pm$ 4.8 [†]	-5.4 $\pm$ 7.0 [†]	58.3 $\pm$ 3.9 [†]	-1.3 $\pm$ 5.2 [†]
RankAlign	neg	84.8 $\pm$ 1.9	57.5 $\pm$ 3.1	79.5 $\pm$ 2.2	55.5 $\pm$ 4.0	43.0 $\pm$ 4.9 [†]	-5.6 $\pm$ 7.0 [†]	53.1 $\pm$ 4.2 [†]	-1.6 $\pm$ 5.6 [†]
Consistency FT	self	85.3 $\pm$ 1.4	62.8 $\pm$ 2.7	—	—	60.4 $\pm$ 3.1 [†]	8.9 $\pm$ 7.5 [†]	—	—
Consistency FT	neg	84.2 $\pm$ 3.3	58.6 $\pm$ 5.0	—	—	46.6 $\pm$ 2.6 [†]	-9.4 $\pm$ 6.8 [†]	—	—
FLORA-PMI	self	92.5 $\pm$ 1.2	77.1 $\pm$ 2.1	81.6 $\pm$ 2.3	60.9 $\pm$ 2.5	80.6 $\pm$ 2.5 [†]	51.5 $\pm$ 4.3 [†]	74.5 $\pm$ 3.5 [†]	36.6 $\pm$ 6.7 [†]
FLORA-Neg	neg	92.3 $\pm$ 1.6	75.9 $\pm$ 3.0	81.2 $\pm$ 2.3	58.8 $\pm$ 3.1	69.6 $\pm$ 3.1 [†]	48.7 $\pm$ 3.3 [†]	68.9 $\pm$ 3.5 [†]	39.5 $\pm$ 5.3 [†]

Table 1: Generator results (ROC_G , ρ) on Hyponymy and IFEval, self-TC and neg-TC eval shown as separate rows (cf. `tab:main-results-multi`, which reports the per-task best). All rows are **epoch 2**. [†] = incomplete (see notes). --- = setting not run for that model (e.g. Consistency FT was not trained for Qwen). **NB:** the Qwen SFT and RankAlign cells were trained *with* `--validator-log-odds` (vlo); the gemma ones were not (effect negligible when $P(\text{Yes})+P(\text{No}) \approx 1$). disc-shots: membership=few, ifeval=zero.

Method	Hyponymy				IFEval			
	G2-9b-it		Q3.5-9b		G2-9b-it		Q3.5-9b	
	ROC_V	Acc_V	ROC_V	Acc_V	ROC_V	Acc_V	ROC_V	Acc_V
Base	94.9 $\pm$ 1.9	87.1 $\pm$ 2.5	95.7 $\pm$ 1.8	88.0 $\pm$ 2.9	78.4 $\pm$ 3.5 [†]	59.7 $\pm$ 3.5 [†]	84.1 $\pm$ 2.8 [†]	70.8 $\pm$ 4.0 [†]
SFT	95.0 $\pm$ 1.8	85.3 $\pm$ 2.5	95.3 $\pm$ 2.0	84.2 $\pm$ 3.4	78.4 $\pm$ 3.7 [†]	66.3 $\pm$ 4.0 [†]	82.8 $\pm$ 3.1 [†]	73.1 $\pm$ 3.8 [†]
RankAlign	95.1 $\pm$ 1.8	85.1 $\pm$ 2.7	95.2 $\pm$ 2.0	87.7 $\pm$ 3.0	52.3 $\pm$ 3.3 [†]	56.2 $\pm$ 3.2 [†]	64.9 $\pm$ 2.6 [†]	50.2 $\pm$ 2.7 [†]
Consistency FT	94.7 $\pm$ 1.9	80.1 $\pm$ 3.0	—	—	78.2 $\pm$ 3.7 [†]	62.9 $\pm$ 4.2 [†]	—	—
FLORA-PMI	94.5 $\pm$ 1.9	84.6 $\pm$ 2.4	95.4 $\pm$ 1.9	82.3 $\pm$ 3.4	79.4 $\pm$ 3.2 [†]	64.5 $\pm$ 3.6 [†]	83.1 $\pm$ 3.3 [†]	73.6 $\pm$ 3.8 [†]
FLORA-Neg	94.6 $\pm$ 2.0	84.9 $\pm$ 2.7	95.2 $\pm$ 2.0	82.5 $\pm$ 3.4	79.0 $\pm$ 3.5 [†]	67.5 $\pm$ 3.9 [†]	83.2 $\pm$ 3.3 [†]	71.8 $\pm$ 3.8 [†]

Table 2: Validator results (ROC_V , Acc_V) on Hyponymy and IFEval. These are independent of the eval-time typicality correction (validator log-odds vs. gold label), so there is no self/neg split. All rows are **epoch 2**.