

RankAlign — RERUN-ONLY results (wandb-rerun checkpoints)

Every value is from the **outputs-rerun-wandb** batch (mll-retrained-for-wandb checkpoints, evaluated 2026-06-07/08/09 in the **qwen35** venv). Cells with no rerun source are ---. The gemma IFEval Base row is blank because no rerun base eval exists for it; the gemma Hyponymy Base row is blank because no rerun base eval exists for gemma-2-9b-it on rosch. All rows epoch 2; RankAlign is *not* epoch-matched to the paper. Each eval-time typicality-correction variant is its own row: **PMI self**/**Neg self** use the model's *own* typicality, **PMI base**/**Neg base** use the *base* model's. A variant row appears only where the rerun batch has it (e.g. Base has own only; FLORA-PMI has no Neg, FLORA-Neg no PMI). disc-shots: membership=few, ifeval=zero.

Generator (ROC_G, ρ)

Method	eval	Hyponymy				IFEval (OOD)				HumanEval (Upper)				HumanEval (Multi)			
		G2-9b-it		Q3.5-9b		G2-9b-it		Q3.5-9b		G4-31b		Q3.5-9b		G4-31b		Q3.5-9b	
		ROC_G	ρ	ROC_G	ρ	ROC_G	ρ	ROC_G	ρ	ROC_G	ρ	ROC_G	ρ	ROC_G	ρ	ROC_G	ρ
Base	PMI self	—	—	73.6 $\pm$ 2.2	45.6 $\pm$ 2.4	—	—	—	—	—	—	69.1 $\pm$ 1.5	30.6 $\pm$ 2.5	—	—	69.0 $\pm$ 1.4	33.8 $\pm$ 2.2
Base	Neg self	—	—	73.8 $\pm$ 2.7	43.9 $\pm$ 3.1	—	—	—	—	—	—	65.7 $\pm$ 1.6	26.3 $\pm$ 2.4	—	—	59.9 $\pm$ 1.6	26.7 $\pm$ 2.2
SFT	PMI self	87.5 $\pm$ 1.4	69.1 $\pm$ 1.9	82.6 $\pm$ 1.9	61.3 $\pm$ 2.8	66.7 $\pm$ 2.8	22.8 $\pm$ 6.0	66.2 $\pm$ 3.6	20.6 $\pm$ 5.6	80.6 $\pm$ 1.2	42.4 $\pm$ 2.6	84.1 $\pm$ 1.2	46.8 $\pm$ 2.3	81.6 $\pm$ 1.1	52.9 $\pm$ 1.8	83.3 $\pm$ 1.1	55.1 $\pm$ 1.8
SFT	PMI base	84.4 $\pm$ 1.4	60.7 $\pm$ 2.2	81.3 $\pm$ 1.8	59.1 $\pm$ 2.5	58.2 $\pm$ 4.0	13.9 $\pm$ 7.3	62.5 $\pm$ 4.1	13.0 $\pm$ 6.2	82.9 $\pm$ 1.2	42.2 $\pm$ 2.9	87.9 $\pm$ 1.1	50.5 $\pm$ 2.2	87.9 $\pm$ 0.9	62.2 $\pm$ 1.7	89.4 $\pm$ 0.9	63.6 $\pm$ 1.6
SFT	Neg self	83.2 $\pm$ 3.3	55.4 $\pm$ 5.2	84.2 $\pm$ 3.0	59.7 $\pm$ 5.3	50.7 $\pm$ 2.7	0.5 $\pm$ 5.2	48.8 $\pm$ 2.8	3.2 $\pm$ 4.5	76.5 $\pm$ 1.5	47.1 $\pm$ 2.5	73.9 $\pm$ 1.6	42.0 $\pm$ 2.5	72.1 $\pm$ 1.5	38.9 $\pm$ 2.5	66.1 $\pm$ 1.5	33.2 $\pm$ 2.4
SFT	Neg base	82.1 $\pm$ 1.3	52.3 $\pm$ 3.0	80.7 $\pm$ 1.9	56.2 $\pm$ 3.4	49.3 $\pm$ 3.8	4.6 $\pm$ 6.7	49.9 $\pm$ 4.3	-3.6 $\pm$ 5.8	81.6 $\pm$ 1.7	40.9 $\pm$ 3.0	94.1 $\pm$ 0.7	54.6 $\pm$ 2.2	93.3 $\pm$ 0.7	68.9 $\pm$ 1.4	94.5 $\pm$ 0.6	67.1 $\pm$ 1.4
RankAlign	PMI self	93.0 $\pm$ 1.5	77.9 $\pm$ 2.2	89.1 $\pm$ 2.0	74.3 $\pm$ 3.0	52.1 $\pm$ 3.9	3.0 $\pm$ 4.5	67.9 $\pm$ 3.6	26.3 $\pm$ 7.6	85.2 $\pm$ 1.1	50.8 $\pm$ 2.5	82.2 $\pm$ 1.3	46.6 $\pm$ 2.1	85.6 $\pm$ 0.9	62.8 $\pm$ 1.6	90.2 $\pm$ 0.8	58.3 $\pm$ 1.7
RankAlign	PMI base	89.5 $\pm$ 1.5	68.8 $\pm$ 2.7	82.6 $\pm$ 2.1	63.2 $\pm$ 3.0	43.4 $\pm$ 4.8	-5.4 $\pm$ 7.0	58.3 $\pm$ 3.9	-1.3 $\pm$ 5.2	86.9 $\pm$ 1.2	48.7 $\pm$ 2.9	87.3 $\pm$ 1.1	53.0 $\pm$ 2.0	90.7 $\pm$ 0.8	68.4 $\pm$ 1.6	93.1 $\pm$ 0.7	62.0 $\pm$ 1.7
RankAlign	Neg self	90.0 $\pm$ 2.4	67.8 $\pm$ 5.2	85.7 $\pm$ 2.6	63.7 $\pm$ 4.7	52.1 $\pm$ 4.7	8.1 $\pm$ 6.7	57.7 $\pm$ 4.2	5.0 $\pm$ 6.8	86.3 $\pm$ 1.3	61.6 $\pm$ 2.3	79.3 $\pm$ 1.4	61.9 $\pm$ 1.9	82.4 $\pm$ 1.5	55.6 $\pm$ 2.7	80.3 $\pm$ 1.4	67.9 $\pm$ 1.7
RankAlign	Neg base	84.8 $\pm$ 2.0	57.7 $\pm$ 3.4	82.3 $\pm$ 2.0	61.9 $\pm$ 3.6	43.0 $\pm$ 4.9	-5.6 $\pm$ 7.0	53.1 $\pm$ 4.2	-1.6 $\pm$ 5.6	88.2 $\pm$ 1.6	54.7 $\pm$ 2.7	90.2 $\pm$ 1.3	60.9 $\pm$ 2.0	94.2 $\pm$ 0.9	75.9 $\pm$ 1.2	96.5 $\pm$ 0.5	65.4 $\pm$ 1.5
Consistency FT	PMI base	—	—	—	—	60.4 $\pm$ 3.1	8.9 $\pm$ 7.5	—	—	83.4 $\pm$ 1.2	48.1 $\pm$ 2.8	86.0 $\pm$ 1.1	46.6 $\pm$ 2.1	87.3 $\pm$ 0.9	58.8 $\pm$ 1.7	88.8 $\pm$ 0.9	57.4 $\pm$ 1.8
Consistency FT	Neg base	—	—	—	—	46.6 $\pm$ 2.6	-9.4 $\pm$ 6.8	—	—	84.4 $\pm$ 1.5	46.2 $\pm$ 3.0	92.6 $\pm$ 0.8	52.7 $\pm$ 2.1	93.3 $\pm$ 0.6	64.0 $\pm$ 1.5	93.9 $\pm$ 0.6	63.5 $\pm$ 1.6
FLORA-PMI	PMI self	93.2 $\pm$ 1.4	77.9 $\pm$ 2.2	86.1 $\pm$ 2.0	69.3 $\pm$ 2.5	84.2 $\pm$ 2.6	53.9 $\pm$ 5.0	79.6 $\pm$ 3.3	46.8 $\pm$ 5.3	84.3 $\pm$ 1.0	57.9 $\pm$ 2.1	78.6 $\pm$ 1.3	56.8 $\pm$ 1.9	84.0 $\pm$ 0.9	64.2 $\pm$ 1.6	82.9 $\pm$ 1.1	59.0 $\pm$ 1.8
FLORA-PMI	PMI base	92.6 $\pm$ 1.3	75.8 $\pm$ 2.1	86.9 $\pm$ 2.0	70.7 $\pm$ 2.1	80.6 $\pm$ 2.5	51.5 $\pm$ 4.3	74.5 $\pm$ 3.5	36.6 $\pm$ 6.7	90.7 $\pm$ 0.8	66.6 $\pm$ 1.8	83.2 $\pm$ 1.2	64.1 $\pm$ 1.7	90.1 $\pm$ 0.8	73.0 $\pm$ 1.4	89.8 $\pm$ 0.9	67.4 $\pm$ 1.6
FLORA-Neg	Neg self	89.1 $\pm$ 2.4	66.8 $\pm$ 5.1	—	—	64.7 $\pm$ 3.6	39.0 $\pm$ 4.3	66.0 $\pm$ 3.0	38.8 $\pm$ 4.7	79.1 $\pm$ 1.5	57.2 $\pm$ 2.3	77.5 $\pm$ 1.5	52.4 $\pm$ 2.1	80.2 $\pm$ 1.4	58.7 $\pm$ 2.3	61.8 $\pm$ 1.7	35.4 $\pm$ 2.6
FLORA-Neg	Neg base	92.0 $\pm$ 1.3	74.5 $\pm$ 3.1	86.7 $\pm$ 2.3	70.2 $\pm$ 3.2	69.6 $\pm$ 3.1	48.7 $\pm$ 3.3	68.9 $\pm$ 3.5	39.5 $\pm$ 5.3	94.8 $\pm$ 0.8	74.9 $\pm$ 1.2	88.9 $\pm$ 1.1	66.8 $\pm$ 1.6	92.0 $\pm$ 0.8	79.3 $\pm$ 1.2	87.9 $\pm$ 1.2	66.8 $\pm$ 1.6

Validator (ROC_V, Acc_V)

Method	Hyponymy				IFEval (OOD)				HumanEval (Upper)				HumanEval (Multi)			
	G2-9b-it		Q3.5-9b		G2-9b-it		Q3.5-9b		G4-31b		Q3.5-9b		G4-31b		Q3.5-9b	
	ROC _V	Acc _V	ROC _V	Acc _V	ROC _V	Acc _V	ROC _V	Acc _V	ROC _V	Acc _V	ROC _V	Acc _V	ROC _V	Acc _V	ROC _V	Acc _V
Base	—	—	95.7 _{±1.8}	88.0 _{±2.9}	—	—	—	—	—	—	81.7 _{±1.3}	65.3 _{±1.4}	—	—	83.4 _{±1.2}	66.8 _{±1.5}
SFT	94.6 _{±2.0}	85.0 _{±2.8}	95.7 _{±1.7}	86.7 _{±3.1}	78.4 _{±3.7}	66.3 _{±4.0}	82.8 _{±3.1}	73.1 _{±3.8}	93.9 _{±0.8}	79.6 _{±1.5}	86.4 _{±1.0}	78.6 _{±1.2}	96.1 _{±0.6}	87.2 _{±1.0}	90.4 _{±0.9}	82.3 _{±1.0}
RankAlign	94.9 _{±1.9}	87.1 _{±2.5}	95.7 _{±1.8}	87.9 _{±3.0}	52.3 _{±3.3}	56.2 _{±3.2}	68.7 _{±4.2}	55.6 _{±3.4}	91.4 _{±1.0}	76.8 _{±1.6}	83.3 _{±1.2}	74.3 _{±1.2}	90.0 _{±1.0}	76.5 _{±1.3}	84.8 _{±1.2}	77.9 _{±1.1}
Consistency FT	—	—	—	—	78.2 _{±3.7}	62.9 _{±4.2}	—	—	93.9 _{±0.8}	87.9 _{±1.0}	84.7 _{±1.1}	76.5 _{±1.2}	93.8 _{±0.8}	83.8 _{±1.3}	88.0 _{±1.1}	80.4 _{±1.1}
FLORA-PMI	94.7 _{±1.8}	83.9 _{±2.7}	95.9 _{±1.7}	86.0 _{±3.2}	79.4 _{±3.2}	64.5 _{±3.6}	83.1 _{±3.3}	73.6 _{±3.8}	93.1 _{±0.9}	84.1 _{±1.4}	86.3 _{±1.1}	76.5 _{±1.3}	94.4 _{±0.7}	82.6 _{±1.3}	89.1 _{±0.9}	80.5 _{±1.1}
FLORA-Neg	94.8 _{±1.9}	86.4 _{±2.6}	95.8 _{±1.7}	86.2 _{±3.0}	79.0 _{±3.5}	67.5 _{±3.9}	83.2 _{±3.3}	71.8 _{±3.8}	94.0 _{±0.8}	88.3 _{±1.1}	86.0 _{±1.1}	78.3 _{±1.1}	94.3 _{±0.8}	76.8 _{±1.3}	88.5 _{±1.0}	80.4 _{±1.1}

Provenance — every value above is from outputs-rerun-wandb/

Cells were aggregated by `_build_ifeval_table.v7.py` / `_build_rosch_table.v7.py` run with `EXTRA_SEARCH_DIR=../outputs-rerun-wandb` (single source, no fallback), then assembled by `_build_rerun_only_table.py`. Checkpoint dirs per column (v6-*=base model, v7-*=trained):

- **Hyponymy / G2-9b-it:** 5 checkpoint(s):

```

- v7-gemma-2-9b-it-d1.43-e2-membership-sans-rosch-v0-all-p0-nv1-ng1-lo0.1-fix1
- v7-gemma-2-9b-it-d2.69-e2-membership-sans-rosch-v0-all-nv1-ng1-vlo-fsx-ppd-sm0.1-fix1
- v7-gemma-2-9b-it-d2.69-e2-membership-sans-rosch-v0-all-tcn-nv1-ng1-vlo-fsx-ppd-sm0.1-fix1
- v7-gemma-2-9b-it-d2.69-e2-membership-sans-rosch-v0-all-tcs-nv1-ng1-vlo-fsx-ppd-sm0.1-fix1
- v7-google-gemma-2-9b-it-delta1.42-epoch2--membership-sans-rosch-v0-all--d2g--random--alpha1.0--full-completion--semi0.1--fix1.merged

```

- **Hyponymy / Q3.5-9b:** 6 checkpoint(s):

```

- v6-Qwen.Qwen3.5-9B
- v7-Qwen3.5-9B-d1.53-e2-membership-sans-rosch-v0-all-p0-nv1-ng1-vlo-lo0.1-fix1
- v7-Qwen3.5-9B-d1.54-e2-membership-sans-rosch-v0-all-nv1-ng1-vlo-fsx-ppd-sm0.1-fix1
- v7-Qwen3.5-9B-d1.54-e2-membership-sans-rosch-v0-all-tcn-nv1-ng1-vlo-fsx-ppd-sm0.1-fix1
- v7-Qwen3.5-9B-d1.54-e2-membership-sans-rosch-v0-all-tcs-nv1-ng1-vlo-fsx-ppd-sm0.1-fix1
- v7-Qwen3.5-9B-d1.54-e2-membership-sans-rosch-v0-all-vlo-sm0.1-fix1

```

- **IFEval / G2-9b-it:** 8 checkpoint(s):

```

- v7-gemma-2-9b-it-d1.29-e2-ifeval-concat-all-p0-nv1-ng1-lo0.1-fix1
- v7-gemma-2-9b-it-d1.38-e2-ifeval-concat-all-sm0.1-fix1
- v7-gemma-2-9b-it-d1.94-e2-ifeval-concat-all-nv1-ng1-vlo-fsx-ppd-sm0.1-fix1
- v7-gemma-2-9b-it-d1.94-e2-ifeval-concat-all-tcn-nv1-ng1-vlo-fsx-ppd-sm0.1-fix1
- v7-gemma-2-9b-it-d1.94-e2-ifeval-concat-all-tcs-nv1-ng1-vlo-fsx-ppd-sm0.1-fix1
- v7-google-gemma-2-9b-it-delta0.98-epoch2--ifeval-concat-all--d2g--random--alpha1.0--full-completion--pref0.0--nllv1.0--nllg1.0--cft--labelonly0.1--fix1.merged
- v7-google-gemma-2-9b-it-delta1.29-epoch2--ifeval-concat-all--d2g--random--alpha1.0--full-completion--pref0.0--nllv1.0--nllg1.0--labelonly0.1--fix1.merged
- v7-google-gemma-2-9b-it-delta1.38-epoch2--ifeval-concat-all--d2g--random--alpha1.0--full-completion--semi0.1--fix1.merged

```

- **IFEval / Q3.5-9b:** 5 checkpoint(s):

```

- v7-Qwen3.5-9B-d0.84-e2-ifeval-concat-all-p0-nv1-ng1-vlo-lo0.1-fix1
- v7-Qwen3.5-9B-d0.96-e2-ifeval-concat-all-nv1-ng1-vlo-fsx-ppd-sm0.1-fix1
- v7-Qwen3.5-9B-d0.96-e2-ifeval-concat-all-tcn-nv1-ng1-vlo-fsx-ppd-sm0.1-fix1
- v7-Qwen3.5-9B-d0.96-e2-ifeval-concat-all-tcs-nv1-ng1-vlo-fsx-ppd-sm0.1-fix1
- v7-Qwen3.5-9B-d0.96-e2-ifeval-concat-all-vlo-sm0.1-fix1

```

- **HumanEval (Upper) & (Multi) / G4-31b & Q3.5-9b:** TEST split, mean $\pm$ SE over 82 problems (per-problem metric, then $\bar{x} \pm \text{sd}/\sqrt{n}$, same estimator as the other columns). ROC_G/ ρ are the *tc* (typicality-corrected) variant; ρ is Spearman. gemma-4-31b has no base eval, so its Base cells are ---. Source score CSVs: `outputs_gemma4_mll_tmp[-multi]/` (gemma-4) and `outputs-rerun-wandb/` (qwen); harvested into `docs/he_harvest.2026-06-17/`. Exact checkpoint dir names per setting: see `docs/he_trainset_eval_model_origins.md`. Note: HumanEval s13 (Consistency FT) also has self-typ evals, not shown here because this shared table only has CFT base-typ rows.